



Towards Performance Portability in GungHo and GOcean

M. Ashworth, **R. Ford**, M. Glover, D. Ham, M.
Hobson, J. Holt, H. Liu, C. Maynard, T. Melvin, L.
Mitchell, E. Mueller, S. Pickles, A. Porter, M.
Rezny, G. Riley, P. Slavin, N. Wood

工合

工合海洋



GungHo

- Globally Uniform, Next Generation, Highly Optimised
- “To research, design and develop a new dynamical core suitable for operational, global and regional, weather and climate simulation on massively parallel computers of the size envisaged over the coming 20 years.”
- Remove the pole problem (replace lat-lon grid)
- Aimed at massively parallel computers – 10^6 way parallel → petaflop
- Split into two phases:
 - 2 years “research” (2011-13)
 - 3 years “development” (2013-2016)
- Met Office, STFC, NERC funding Universities of: Bath, Exeter, Imperial, Leeds, Manchester, Reading

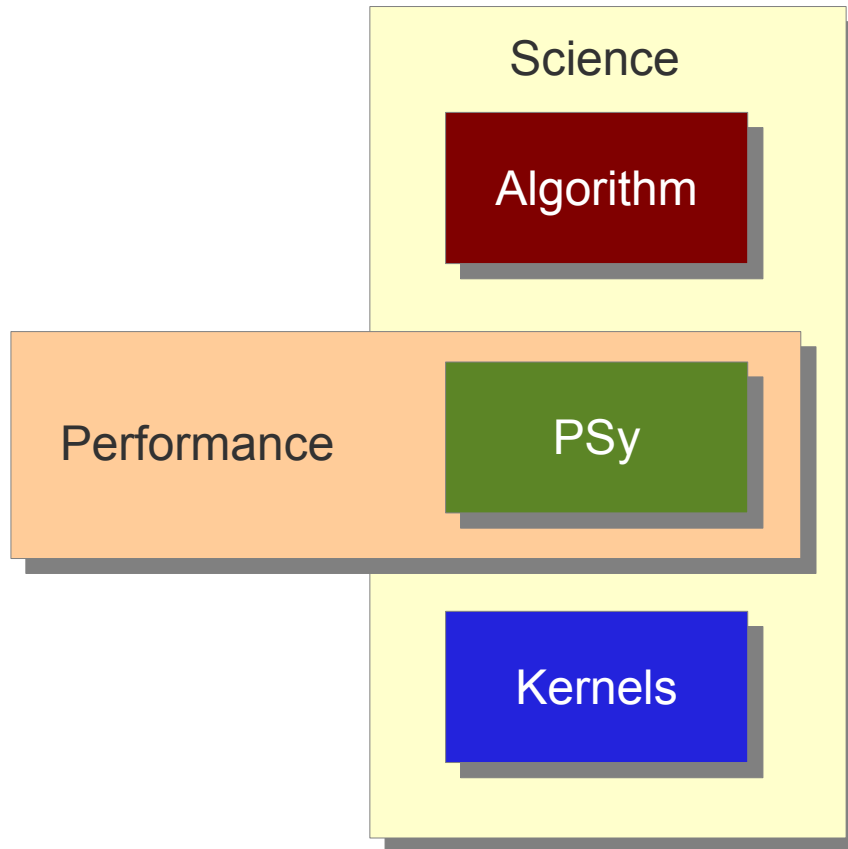


GungHo Status

- (Most likely) Cubed Sphere
 - Extruded (columnar) mesh (2d+1d)
- Multi-grid
- Finite element approach
- Dynamo implementation



PSyKAI



GH Algorithm

- Hand written (conforming to Fortran 2003)
- Use Field objects
- Logically operate on global fields
- Invoke approach
 - Algorithm layer **specifies** what the PSy layer has to do
 - Algorithm layer “specifications” will be pre-processed to specific calls which replace original
 - Invocation can take a 'list' of kernel specs


```
call invoke(kern(field1,field2,field3)...)

```



Illustration : Alg

```
...  
call invoke(&  
    rhs_v3_type(rhs)&  
    )  
...
```




```
...  
use psy, only: invoke_rhs_v3  
...  
call invoke_rhs_v3(rhs)  
...
```



Multi-function Illustration : Alg

```
...  
call invoke(&  
  set(res_norm, 0.0), &  
  galerkin_action(x, Mu, u), &  
  galerkin_matrix_free_update(u, Mu, b, M_l, res_norm) &  
  )  
...
```



```
...  
USE psy, ONLY: invoke_2  
...  
CALL invoke_2(b, m_l, mu, u, x, res_norm)  
...
```



GH Kernel

- Hand written conforming to the GungHo PSyKAI API
- Scientific
 - There will also be library routines e.g. linear algebra
- Column based (written assuming k-inner)
- Access raw data (not field objects)
- Associated metadata (required for code generation) e.g.
 - Intents (extending fortran's in and out)
 - The function space a field is on (v0, v1, v2, ...)
 - what the kernel iterates over (cells, edges, ...)



Illustration GH0.1 API

```
module rhs_v3_mod
...
type, public, extends(kernel_type) :: rhs_v3_type
  private
  type(arg_type) :: meta_args(1) = [ &
    arg_type(gh_rw,v3,fe,.true.,.false.,.true.) &
  ]
  integer :: iterates_over = cells
contains
  procedure, nopass :: rhs_v3_code
end type
...
subroutine rhs_v3_code(nlayers,ndf,map,v3_basis,x,gq)
...
end subroutine rhs_v3_code
end module rhs_v3_mod
```



GH PSy

- The Optimised PSy may be generated
 - Manual “reference/vanilla” version
 - Should be easily debuggable
- Functional responsibility
 - iterating over columns
 - Mapping of algorithm fields types/objects to data required by kernel
 - Number of arguments may not be the same (e.g. dof information)
 - Halo exchange
- Performance responsibility
 - Optimise for particular architectures → portable performance
 - Threading: OpenMP, OpenACC, ..., Kernel re-ordering, Fusion, Inlining, ...



Illustration GH0.1 API

```
module psy
...
subroutine invoke_rhs_v3(rhs)
  use rhs_v3_mod, only : rhs_v3_code
  ...
  nlayers=rhs%get_nlayers()
  ndf = rhs%vspace%get_ndf()
  call rhs%vspace%get_basis(v3_basis)
  do cell = 1, rhs%get_ncell()
    call rhs%vspace%get_cell_dofmap(cell,map)
    call rhs_v3_code(nlayers,ndf,map,v3_basis,rhs%data,rhs%gaussian_quadrature)
  end do
end subroutine invoke_rhs_v3
...
end module psy
```

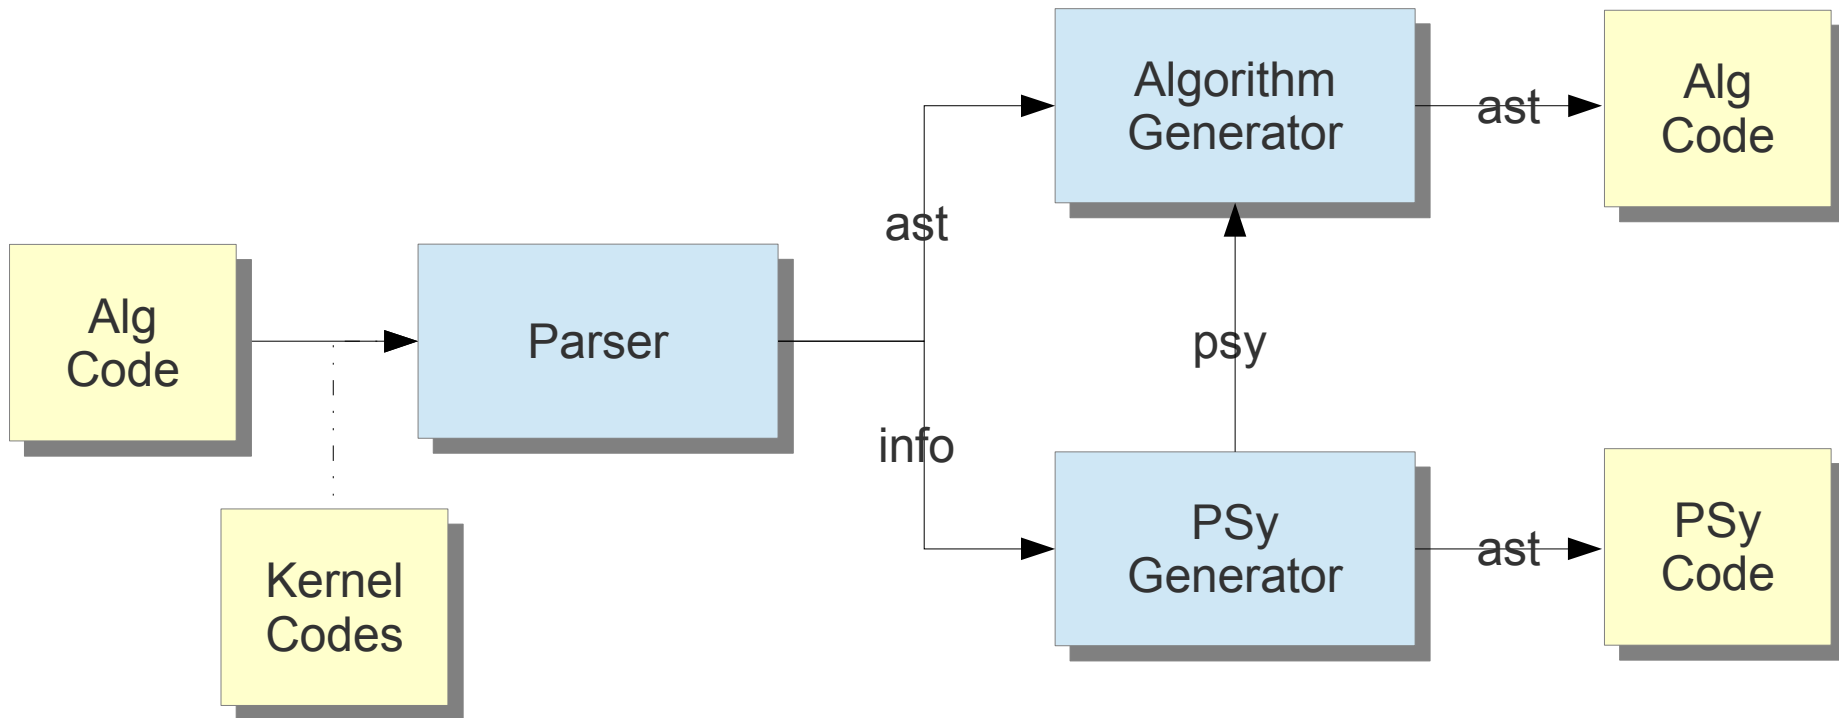


PSyclone Code Generation

- Code generation requires, Alg API, Kern API, Kernel metadata
- Code generation can help with
 - optimisation – labourious and error prone by hand
 - changes in interfaces
- PSyclone:
 - Taking an interactive optimisation approach to support the expert
 - Could also offer full automation option at a later date
 - Generates correct sequential code for GH 0.1 API
 - 4,113 lines of Python code
 - Following optimisations are available:
 - Loop fusion
 - OMP loop parallelisation
 - Loop colouring
 - Kernel inlining



PSyclone



PSyclone


```
> python generate.py -oalg alg.f90 -opsy psy.f90 -api dynamo0.1 example.f90
```

```
>>> from generator import generate  
>>> psy, alg = generate("example.f90", api="dynamo0.1")  
>>> print str(psy.gen)  
>>> print str(alg.gen)
```

```
>>> from algGen import Alg  
>>> from parser import parse  
>>> from psyGen import PSyFactory  
>>> ast, info = parse("example.f90", api="dynamo0.1")  
>>> psy = PSyFactory("dynamo0.1").create(info)  
>>> alg = Alg(ast,psy)  
>>> print str(psy.gen)  
>>> print str(alg.gen)
```



PSy Schedule




```

>>> psy = PSyFactory("dynamo0.1").create(info)
>>> invokes = psy.invokes
>>> invokes.names
>>> invoke = invokes.get("name")
>>> schedule = invoke.schedule
>>> schedule.view()
  
```

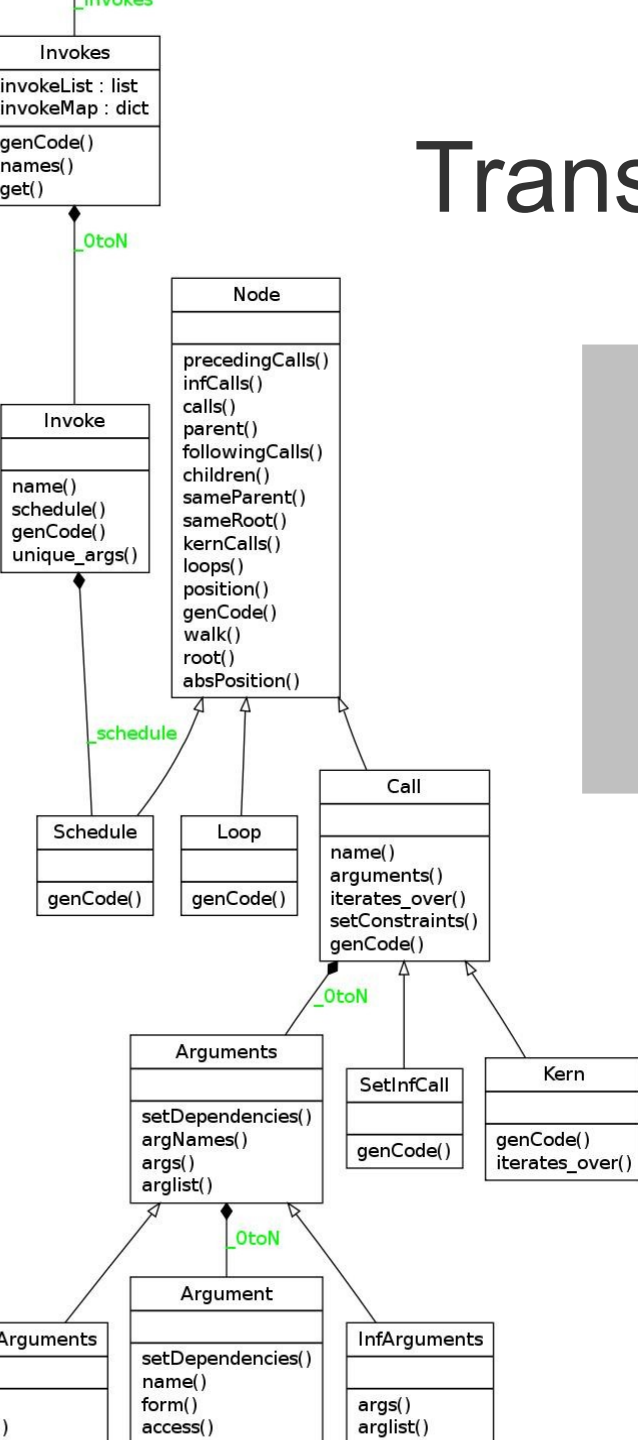


Schedule Illustration

```
...  
call invoke(&  
    set(res_norm, 0.0), &  
    galerkin_action(x, Mu, u), &  
    galerkin_matrix_free_update(u, Mu, b, M_l, res_norm)&  
    )  
...
```




Transform PSy Schedule



```

>>> lf = LoopFuseTrans()
>>> loop1 = schedule.children[0]
>>> loop2 = schedule.children[1]
>>> new_schedule, memento = lf.apply(loop1, loop2)
>>> invoke._schedule = new_schedule
    
```



GOcean

- NERC funded Proof-of-principle 1 year project
- STFC & NOC + advice from GungHo colleagues
- Can GungHo PSyKAI approach be applied to Ocean Models?
- Finite Difference
- 2+1 test codes: shallow, NEMO-lite-2D, NEMO-lite-3D



GOcean

- T, P, U, V metadata to describe staggering
- API: Point-wise kernels, direct addressing
- Manual shallow and NEMO-lite-2D using GOcean PSyKAI
- PSyclone correct sequential code for 0.1 API
- Loop fusion, OpenMP loop parallel, inlining supported
- Nearly 90% of PSyclone is common




Original Shallow

```
...  
DO J=1,N  
  DO I=1,M  
    CU(I+1,J) = .5*(P(I+1,J)+P(I,J))*U(I+1,J)  
    CV(I,J+1) = .5*(P(I,J+1)+P(I,J))*V(I,J+1)  
    Z(I+1,J+1) =(FSDX*(V(I+1,J+1)-V(I,J+1))-FSDY*(U(I+1,J+1) &  
      -U(I+1,J)))/(P(I,J)+P(I+1,J)+P(I+1,J+1)+P(I,J+1))  
    H(I,J) = P(I,J)+.25*(U(I+1,J)*U(I+1,J)+U(I,J)*U(I,J)   &  
      +V(I,J+1)*V(I,J+1)+V(I,J)*V(I,J))  
  END DO  
END DO  
...
```



Shallow Alg

```
...  
call invoke( compute_cu_type(CU, P, U), &  
             compute_cv_type(CV, P, V), &  
             compute_z_type(Z, P, U, V), &  
             compute_h_type(H, P, U, V) )  
...
```



```
...  
USE psy_shallow, ONLY: invoke_0  
...  
CALL invoke_0(cu, p, u, cv, v, z, h)  
...
```



Shallow Kern

```
module compute_cu_mod
  use kind_params_mod
  ...
  type, extends(kernel_type) :: compute_cu_type
    type(arg), dimension(3) :: meta_args = &
      (/ arg(WRITE, CU, POINTWISE), & ! cu
        arg(READ, CT, POINTWISE), & ! p
        arg(READ, CU, POINTWISE) & ! u
      /)
    integer :: ITERATES_OVER = DOFS
  contains
    procedure, nopass :: code => compute_cu_code
  end type compute_cu_type
  ...
  subroutine compute_cu_code(i, j, cu, p, u)
    ...
    CU(I,J) = .5*(P(I,J)+P(I-1,J))*U(I,J)
  end subroutine compute_cu_code
end module compute_cu_mod
```



Example

```
SUBROUTINE invoke_0(cu_1, p, u, cv_1, v, z, h)
...
DO j=cu%jstart,cu%jstop
  DO i=cu%istart,cu%istop
    CALL compute_cu_code(i, j, cu_1, p, u)
  END DO
END DO
DO j=cv%jstart,cv%jstop
  DO i=cv%istart,cv%istop
    CALL compute_cv_code(i, j, cv_1, p, v)
  END DO
END DO
...
END SUBROUTINE invoke_0
```



GOcean shallow 128

Compiler:	Cray 8.3.3	Intel 14.0.1	Gnu 4.8.2	Intel 14.0.0
Hardware:	IvyBridge	IvyBridge	Haswell	Haswell
Original	0.29	0.40	0.37	0.37
Vanilla	0.41	0.49	6.30	0.42
Explicit bounds	0.34	0.47	6.34	0.43
In-lined kernels	0.35	0.47	0.55	0.42
Loop fused	0.34	0.43	0.53	0.39
In-lined copy	0.34	0.43	0.54	0.39
Fused copy	0.31	0.51	0.54	0.45
Fastest	0.31	0.43	0.53	0.39
% slower	4.26	7.30	42.25	5.43



GOcean shallow sizes with Cray compiler

Problem size	64	128	256	512	1024
Original	0.008	0.29	1.21	5.70	44.12
Fastest	0.008	0.31	1.3	5.88	42.77
% slower	-.23	4.26	7.78	3.18	-3.06




Summary

- PSyKAI approach shows promise
- Code Generation shows promise
- Initial results show promise
- Can promise become practice?

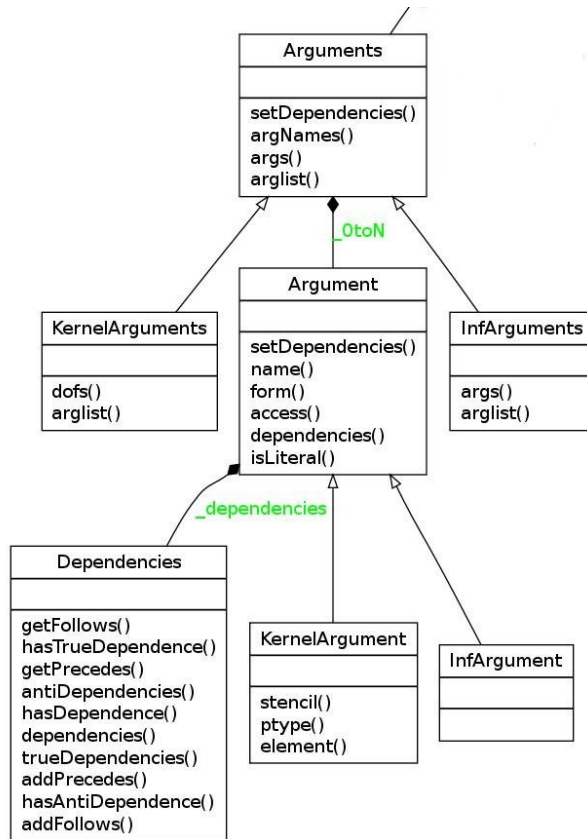


Gocean shallow

Performance Relative to Original



Dependencies



set(res_norm, 0.0)

galerkin_action(x, Mu, u)

galerkin_matrix_free_update(u, Mu, b, M_l, res_norm)

