



# 工合海洋 GOcean

An update...

**Andrew Porter**, Rupert Ford & Graham Riley  
STFC

Hedong Liu & Jason Holt  
NOC

# Project Overview

- NERC funded, 03/2014 – 02/2015
- Collaboration between NOC and STFC
- Existing NEMO code structure needs updating
  - 20 yrs old and MPI only, vertical level index last
- Investigate the feasibility of applying the GungHo approach to ocean modelling
- Ocean models can avoid pole problem
  - put poles over land; can retain a lat-lon grid
- **Look to extend the developing Gung-Ho infrastructure to support finite difference on lat-lon grids**

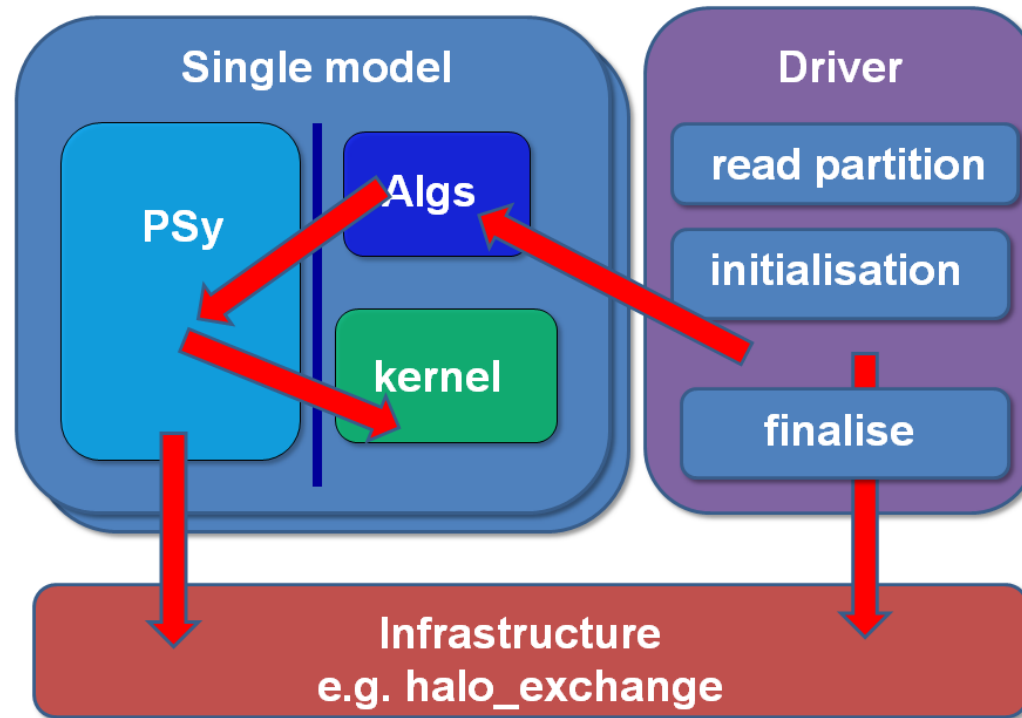


# The Plan...

- Obtain finite-difference, shallow-water model(s)
- Re-structure following PSyKAI approach
- Extend kernel meta-data to capture necessary information
  - Finite difference, lat-lon, direct-addressed
- Extend PSyclone to generate the Parallel System (middle layer)
  - Support for multi/many-core/GPU etc. but not MPI
- Investigate performance implications and feedback to PSyclone
- Feedback to the NEMO system team



# Separation of Concerns in Gung-Ho (recap)



# The Parallel System, Kernel, Algorithm (PSyKAI) Approach...

- Oceanographer writes the algorithm and kernel layers, following certain rules
  - no need to worry about relative indexing of various fields
  - no need to worry about parallelism
- A code-generation system (PSyclone) generates the PSy middle layer
  - glues the algorithm and kernels together
  - incorporates all code related to parallelism



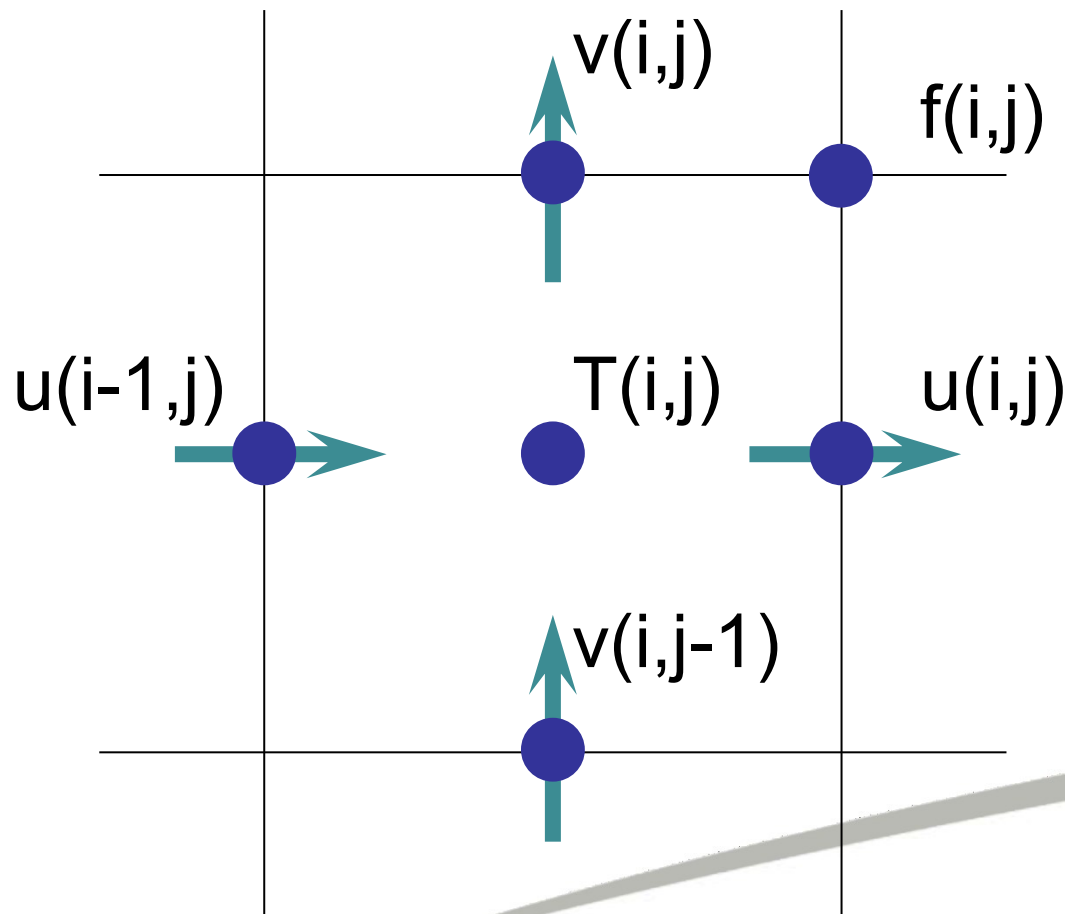
# Two shallow-water codes...

- We are applying PSyKAI approach to two codes:
  - ‘shallow:’ originally written by Swarztrauber, NCAR
  - ‘NEMOLite2D:’ 2D, free surface part of NEMO extracted by NOC
- Both use Finite Difference on Arakawa C grid
- But there are important differences:
  - Boundary conditions (periodic vs. forced/closed)
  - Relative indexing of variables on the grid
- Understanding and expressing these differences is essential for correct code generation



# Placing variables...

- Variable placement for the Staggered Arakawa C-grid with NEMO indexing convention:



$T$ : scalars (density, salinity etc.)

$u, v$ : velocity components

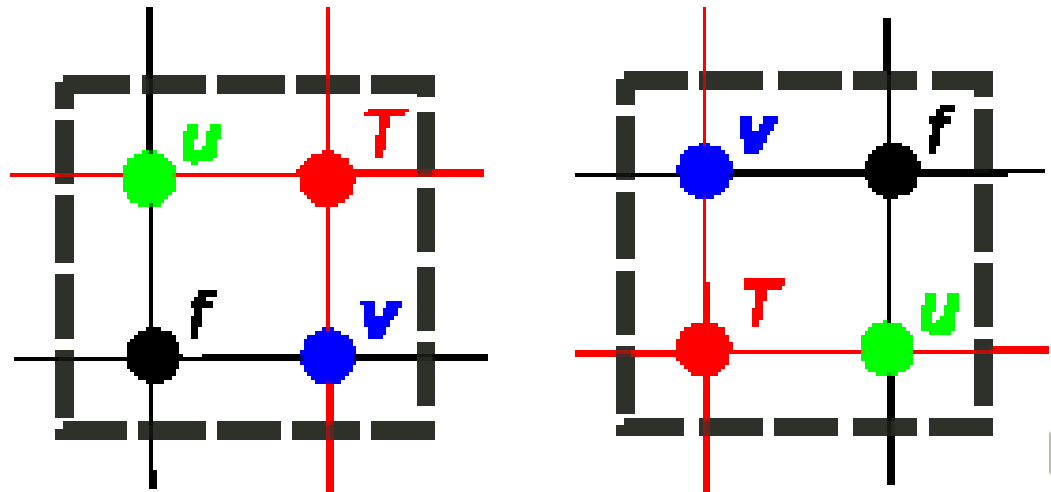
$f$ : vorticity



# Offset choice

- A *developer* can choose how the different grid-point types are indexed relative to T
- **shallow** defines  $\{u, v, f\}$  points to the South and West of the T point to have same  $(i, j)$  index while **NEMO** uses those to the North and East:

- We call this choice the ‘**offset**’ of the grids
- Specified in kernel meta-data





# PSyKAI-fication...

Take shallow-water code, e.g.:

```
DO ncycle=1,itmax    ! Time-stepping loop
! COMPUTE CAPITAL U, CAPITAL V, Z AND H
  DO J=1,N
    DO I=1,M
      CU(I+1,J) = &
        .5*(P(I+1,J)+P(I,J))*U(I+1,J)
      CV(I,J+1) = &
        .5*(P(I,J+1)+P(I,J))*V(I,J+1)
      ...
    END DO
  END DO
```



# Construct (point-wise) kernels, e.g.:

```
SUBROUTINE compute_cu_code(i, j, &  
                           cu, p, u)  
  
  INTEGER,    INTENT(in)  :: i, j  
  REAL(wp),   INTENT(in)  :: p(:, :), u(:, :)  
  REAL(wp),   INTENT(inout):: cu(:, :)  
  
  CU(I, J) = .5* (P(I, J) + P(I-1, J)) * U(I, J)  
  
END SUBROUTINE compute_cu_code
```



**Re-structure following Gung-Ho rules (all computation must be done in a kernel):**

**Time-stepping loop at algorithm level becomes....**

```
DO ncycle=1,itmax ! Time-stepping loop
```

```
! COMPUTE CAPITAL U, CAPITAL V, Z AND H
```

```
CALL compute_cu(CU, P, U)
```

```
CALL compute_cv(CV, P, V)
```

```
...
```

```
CALL invoke (compute_cu(CU, P, V), &  
             compute_cv(CV, P, V), &  
             ...)
```



# Translation/generation...

```
PROGRAM shallow
  USE psy_shallow, ONLY: invoke_0
  ...
  ! ** Start of time loop **
  DO ncycle=1,itmax

    ! COMPUTE CAPITAL U, CAPITAL V, Z, H
    CALL invoke_0(cu, p, u, cv, v, z, h)
```



```
SUBROUTINE invoke_0(cu_1,p,u,cv_1,v,z,h)
  USE compute_cv_mod, ONLY: compute_cv_code
  USE compute_cu_mod, ONLY: compute_cu_code
  USE topology_mod,    ONLY: cu, cv
  REAL,intent(inout),dimension(:, :) :: cu_1,p,u,cv_1,v,z,h
  DO i=cu%istart,cu%istop
    DO j=cu%jstart,cu%jstop
      CALL compute_cu_code(i, j, cu_1, p, u)
    END DO
  END DO
  DO i=cv%istart,cv%istop
    DO j=cv%jstart,cv%jstop
      CALL compute_cv_code(i, j, cv_1, p, v)
    ...
```

# Kernel meta-data

- Kernels make use of several grid-related quantities, e.g. area of cell around a T point, T-point mask etc.

```
subroutine next_sshu_code(ji,jj, sshn_u, sshn, &
                        tmask,e12t,e12u)
  implicit none
  integer,          intent(in)      :: ji, jj
  integer, dimension(:,,:), intent(in) :: tmask
  real(wp), dimension(:,,:), intent(in)  :: e12t, e12u
  real(wp), dimension(:,,:), intent(inout) :: sshn_u
  _ real(wp), dimension(:,,:), intent(in)  :: sshn
```

- The algorithm layer should not/cannot supply these:

```
call invoke( next_sshu(sshn_u, sshn), &
             next_sshv(sshn_v, sshn) )
```



- Extend meta-data to specify what quantities a kernel requires from the infrastructure:

```
type, extends(kernel_type) :: next_sshu
    type(arg), dimension(5) :: meta_args = &
        (/ arg(READWRITE, CU, POINTWISE), &
           arg(READ, CU, POINTWISE), &
           arg(READ, GRID_MASK_T), &
           arg(READ, GRID_AREA_T), &
           arg(READ, GRID_AREA_U) &
        /)
```

- PSyclone must then supply these quantities from the generated middle layer (e.g. by accessing a grid object pointed to by a field object)



# What about performance?

Some timings for 2000 time steps of 128x128 domain on a single core:

Compiler:	Gnu 4.8.2	Intel 14.0.0	Cray 8.2.6	Intel 14.0.1
CPU:	Xeon E5-1620 v3, 3.7 GHz (Haswell)	Xeon E5-1620 v3, 3.7 GHz (Haswell)	Xeon E5-2697, 2.7 GHz (Ivy Bridge)	Xeon E5-2697, 2.7 GHz (Ivy Bridge)
Original*:	0.37	0.37	0.29	0.40
Manual, vanilla†:	6.30	0.42	0.41	0.49
Manual, opt'd‡:	0.53	0.39	0.31	0.43

\* Unmodified shallow code

† Manual PSyKAI version

‡ As † but with middle layer optimised

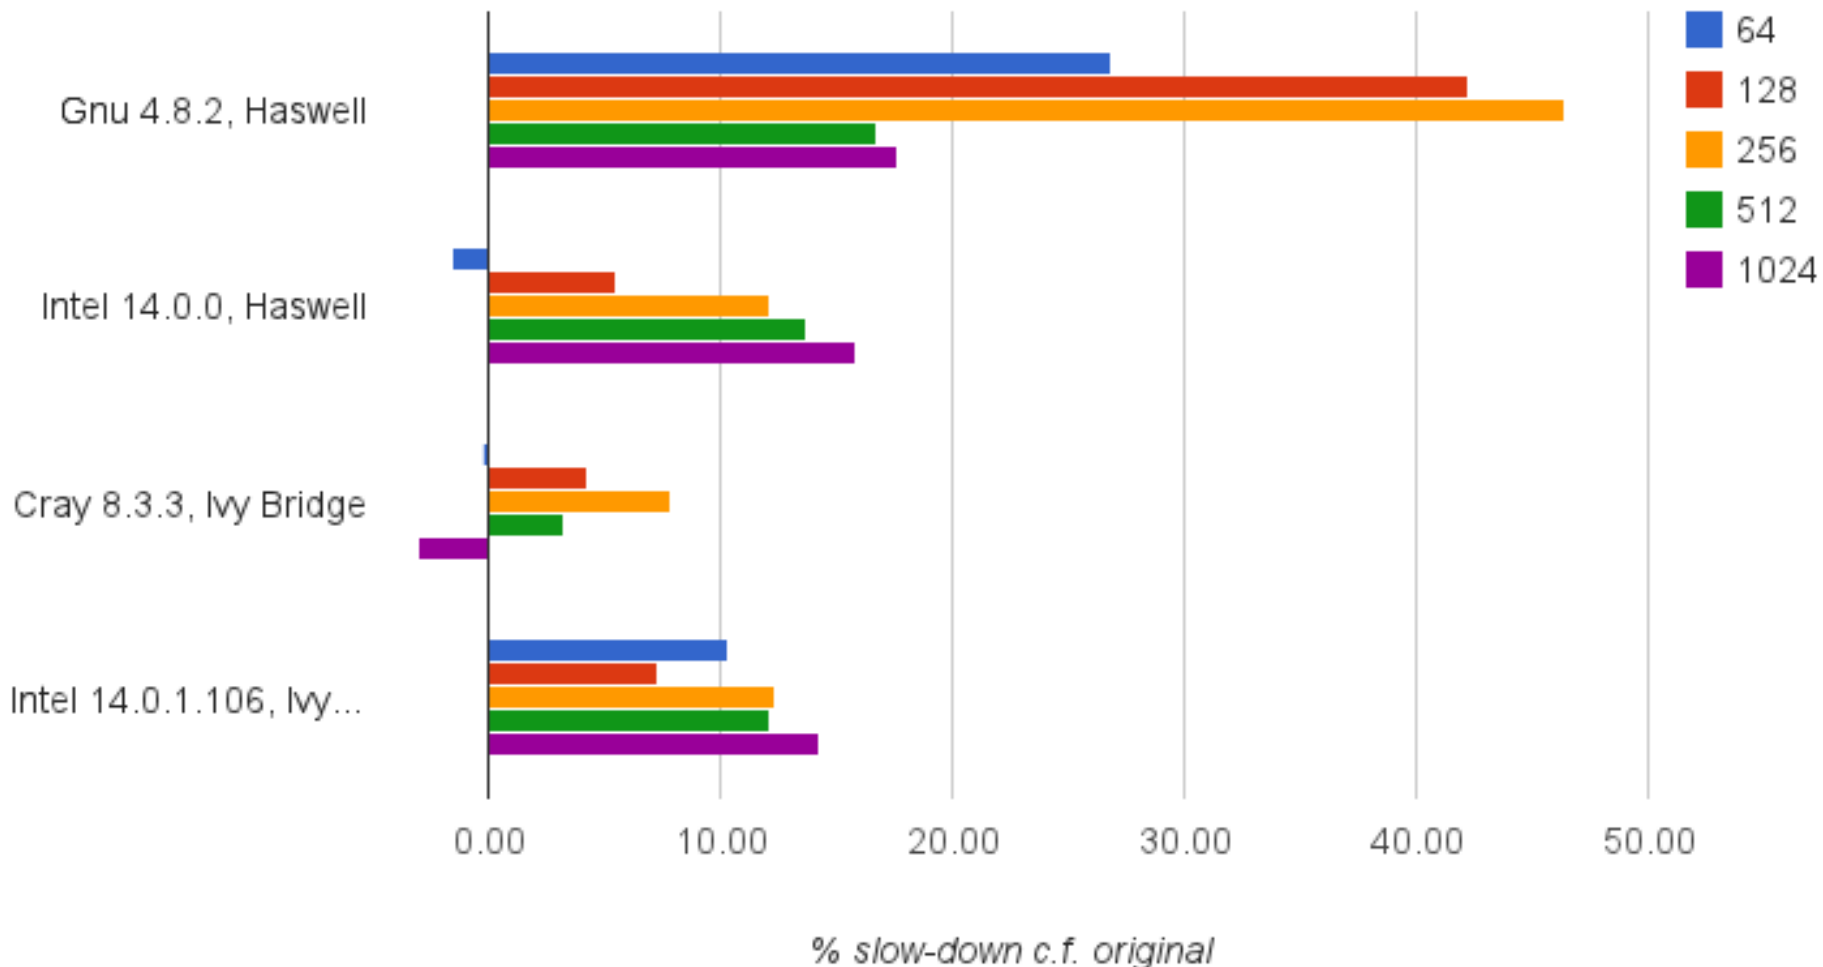


Science & Technology  
Facilities Council



# All compilers are not created equal

Performance of best (so far) PSyKAI version compared to original, unmodified 'shallow':



# Next steps...

- Supply vendors (IBM, NVIDIA) with original and PSyKAI'ised versions and let them optimise
  - Use lessons learned to continue to improve PSyclone
- Gain fuller understanding of cost (if any) of introducing layered structure
- Currently not exploring the optimisation space
  - Only attempting to recover original code structure
- Three dimensions
  - Current test cases are two-dimensional
  - Full models are a mixture of 2D and 3D...
  - NOC working on introducing some 3D aspects to NEMOLite



# Summary I

- Separation of Concerns: A practical approach to marrying the requirements of oceanographers with the requirements for performance in the (pre-) exascale era.
- Introduces flexibility needed to achieve performance on different architectures
  - potentially enables e.g. OpenMP or OpenACC to be used, depending on target hardware
- Very dependent on middle layer to retrieve the performance that we've thrown out



# Summary II

- Framework now supports two distinct shallow-water models
- Basic code generation functional
  - Support for loop fusion and OpenMP transformations
- Working in collaboration with hardware vendors to understand what loses us performance and how to regain it
- Potential to explore optimisation space that has been opened by the flexibility provided by code transformation & generation



# Extras...



**Science & Technology**  
Facilities Council

# Middle layer is generated...

```
!  ** Start of time loop **
```

```
DO ncycle=1,itmax
```

```
! COMPUTE CAPITAL U, CAPITAL V, Z, H
```

```
call invoke(compute_cu_type(CU, P, U), &  
            compute_cv_type(CV, P, V), &  
            compute_z_type(z, P, U, V), &  
            compute_h_type(h, P, U, V) )
```



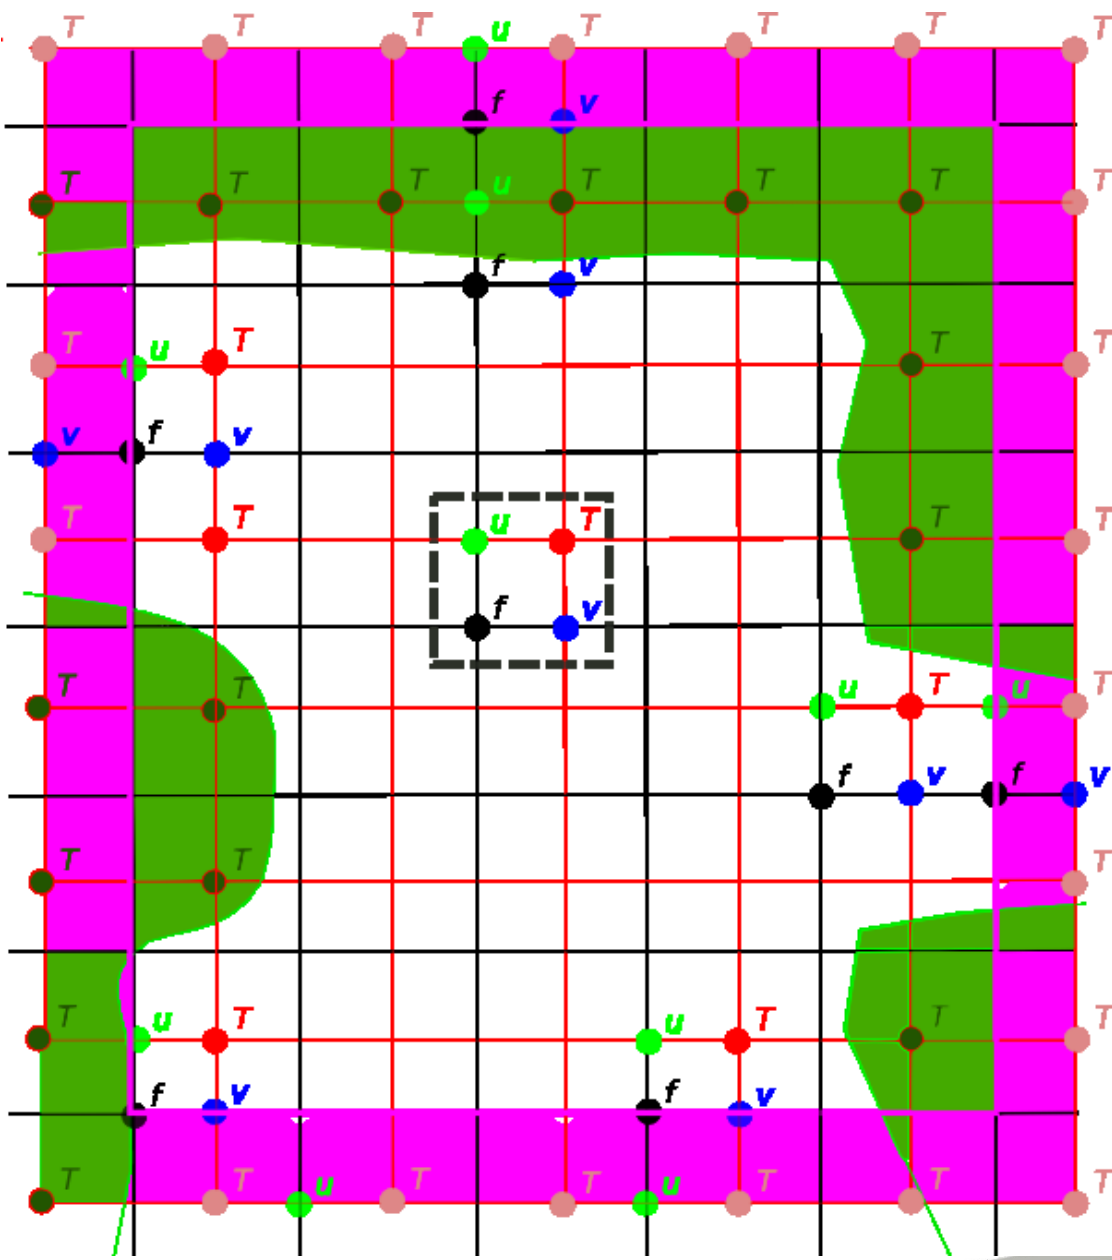
## Data movement for PBC update

## Boundary region

## External/boundary grid points

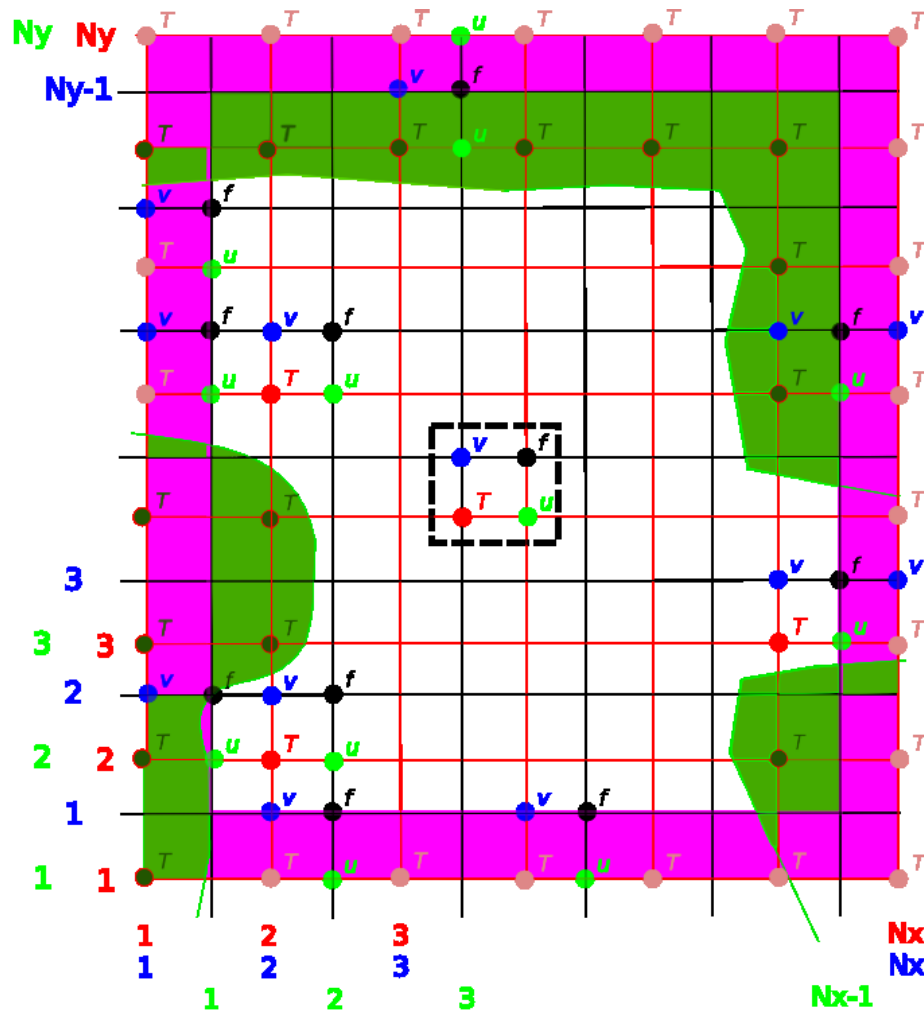
# SW offset with open and closed boundaries

- User defines domain in terms of T points
- Definition *includes* the boundary points
  - Serial implementation doesn't need halos





# NE offset with in-place boundary conditions



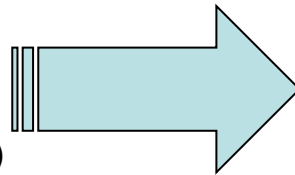
- Model domain
- Boundary region
- T-point outside domain
- V-point inside domain
- T-point inside domain
- Boundary V-point
- Land point
- Grid pts with same (i,j)



# Code generation to the rescue...

- PSyclone currently has rudimentary support for loop-fusion (and the addition of OpenMP)
- e.g. for the time-smoothing section of the code where all 3 loops have same bounds:

```
DO j=1, SIZE(uold, 2)
  DO i=1, SIZE(uold, 1)
    CALL tsmooth_code(i, j, u...)
  DO i=1, SIZE(vold, 1)
    DO j=1, SIZE(vold, 2)
      CALL tsmooth_code(i, j, v...)
    DO i=1, SIZE(pold, 1)
      DO j=1, SIZE(pold, 2)
        CALL tsmooth_code(i, j, p...)
```



```
DO j=1, SIZE(uold, 2)
  DO i=1, SIZE(uold, 1)
    CALL tsmooth_code(i, j, u...)
    CALL tsmooth_code(i, j, v...)
    CALL tsmooth_code(i, j, p...)
  END DO
END DO
```

